

Word Count With Two Text file

Code

```
package two_file_wc;

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class two_file_wc {

    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(LongWritable key, Text value, Context context)
            throws IOException, InterruptedException {
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
        public void reduce(Text key, Iterable<IntWritable> values, Context context)
            throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            context.write(key, new IntWritable(sum));
        }
    }
}
```

```
        }
    }
}

public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

public static void main(String[] args) throws Exception {

    if (args.length != 3) {
        System.err.println("Usage: WordCountTwoFiles <input path1> <input path2> <output
path>");
        System.exit(-1);
    }

    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "Word Count for Two Files");

    job.setJarByClass(two_file_wc.class);
    job.setMapperClass(Map.class);
    job.setReducerClass(Reduce.class);

    job.setMapOutputKeyClass(Text.class);
    job.setMapOutputValueClass(IntWritable.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
}
```

```

job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);

// Add both input files
FileInputFormat.addInputPath(job, new Path(args[0]));
FileInputFormat.addInputPath(job, new Path(args[1]));

Path outputPath = new Path(args[2]);
FileOutputFormat.setOutputPath(job, outputPath);

// Delete output folder if it exists
outputPath.getFileSystem(conf).delete(outputPath, true);

System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Code with explanation

package two_file_wc;

This defines the package name where your Java class belongs. It helps in organizing your Hadoop programs.

import java.io.IOException;

Imports the **IOException** class to handle input/output-related errors.

import java.util.StringTokenizer;

Imports **StringTokenizer** — used to split a line of text into individual words (tokens).

import org.apache.hadoop.conf.Configuration;

Used to access Hadoop configuration settings.

import org.apache.hadoop.fs.Path;

Represents a file or directory path in the Hadoop File System (HDFS).

```
import org.apache.hadoop.io.IntWritable;
```

A Hadoop data type for storing integer values (used instead of Java's `int`).

```
import org.apache.hadoop.io.LongWritable;
```

A Hadoop data type for storing long integer values.

```
import org.apache.hadoop.io.Text;
```

A Hadoop data type for strings (used instead of Java's `String`).

```
import org.apache.hadoop.mapreduce.Job;
```

Represents a MapReduce job configuration (including Mapper, Reducer, input, output paths, etc.).

```
import org.apache.hadoop.mapreduce.Mapper;
```

Base class for writing the Mapper logic.

```
import org.apache.hadoop.mapreduce.Reducer;
```

Base class for writing the Reducer logic.

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

Used to define the input file(s) or directory for the MapReduce job.

```
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
```

Specifies that the input will be plain text files, where each line is a record.

```
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

Used to define the output directory for the MapReduce job.

```
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

Specifies that the output will be stored as plain text.

```
public class two_file_wc {
```

This defines your main public class. It contains both the Mapper, Reducer, and main method.

Mapper Class

```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {
```

Defines an inner static class named `Map` that extends Hadoop's `Mapper` class.

It takes input key-value pairs:

- Key → `LongWritable` (line offset)
- Value → `Text` (line content)
And outputs key-value pairs:
 - Key → `Text` (a word)
 - Value → `IntWritable` (count 1)

```
private final static IntWritable one = new IntWritable(1);
```

Creates a constant value `1`, which represents one occurrence of a word.

```
private Text word = new Text();
```

Creates a reusable `Text` object to store each word.

```
public void map(LongWritable key, Text value, Context context) throws IOException,  
InterruptedException {
```

This is the Mapper's main method — Hadoop calls it once for each line of input.

```
String line = value.toString();
```

Converts the line from `Text` to a regular Java `String`.

```
StringTokenizer tokenizer = new StringTokenizer(line);
```

Splits the line into individual words using spaces as delimiters.

```
while (tokenizer.hasMoreTokens()) {
```

Loops through all the words in the line.

```
word.set(tokenizer.nextToken());
```

Stores the next word in the `word` variable.

```
context.write(word, one);
```

Emits the word along with the count `1` to the Hadoop framework (Mapper output).

Reducer Class

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
```

Defines a `Reduce` class that extends Hadoop's `Reducer`.

It takes input from Mapper:

- Key → `Text` (word)
 - Values → list of `IntWritable` (counts)
And outputs:
 - Key → `Text` (word)
 - Value → `IntWritable` (final total count)
-

```
public void reduce(Text key, Iterable<IntWritable> values, Context context) throws  
IOException, InterruptedException {
```

This method is called once for each unique word key.

```
int sum = 0;
```

Initializes a counter for total occurrences of a word.

```
for (IntWritable val : values) { sum += val.get(); }
```

Iterates over all counts for that word and adds them up.

```
context.write(key, new IntWritable(sum));
```

Writes the final word and its total count to output.

Main Method

```
public static void main(String[] args) throws Exception {
```

Main function — entry point of the program.

```
if (args.length != 3) { ... }
```

Checks whether exactly three arguments are given (two input files and one output directory). If not, it prints a usage message and exits.

```
Configuration conf = new Configuration();
```

Creates a Hadoop configuration object to hold job settings.

```
Job job = Job.getInstance(conf, "Word Count for Two Files");
```

Creates and names a new MapReduce job instance.

```
job.setJarByClass(two_file_wc.class);
```

Specifies which class contains the job's main method (used to locate your JAR).

```
job.setMapperClass(Map.class);
```

Tells Hadoop which Mapper class to use.

```
job.setReducerClass(Reduce.class);
```

Tells Hadoop which Reducer class to use.

```
job.setMapOutputKeyClass(Text.class);
```

Defines the Mapper output key type (word).

```
job.setMapOutputValueClass(IntWritable.class);
```

Defines the Mapper output value type (count 1).

```
job.setOutputKeyClass(Text.class);
```

Defines the final output key type (word).

```
job.setOutputValueClass(IntWritable.class);
```

Defines the final output value type (total count).

```
job.setInputFormatClass(TextInputFormat.class);
```

Specifies that the input files are text files.

```
job.setOutputFormatClass(TextOutputFormat.class);
```

Specifies that the output will be text format.

```
FileInputFormat.addInputPath(job, new Path(args[0]));
```

Sets the first input file path.

```
FileInputFormat.addInputPath(job, new Path(args[1]));
```

Adds the second input file path. Both files will be processed together.

```
Path outputPath = new Path(args[2]);
```

Specifies the output directory path.

```
FileOutputFormat.setOutputPath(job, outputPath);
```

Sets where the result will be stored in HDFS.

```
outputPath.getFileSystem(conf).delete(outputPath, true);
```

Deletes the output directory if it already exists to prevent job failure.

System.exit(job.waitForCompletion(true) ? 0 : 1);

Runs the job and exits:

- Returns **0** if successful
- Returns **1** if it fails
-

- : Commands :-

```
administrator@administrator:~$ hadoop jar  
/home/administrator/HadoopJar/TwoFileWC.jar two_file_wc.two_file_wc  
/data/TFWC/T1 /data/TFWC/T2 /data/TFWCoutput
```

- administrator@administrator:~\$ hadoop fs -ls /data/
Found 2 items

```
drwxr-xr-x    - administrator supergroup          0 2025-10-07 16:08 /data/TFWC  
drwxr-xr-x    - administrator supergroup          0 2025-10-07 16:12  
/data/TFWCoutput
```

- administrator@administrator:~\$ hadoop fs -ls /data/TFWCoutput
Found 2 items

```
-rw-r--r--    1 administrator supergroup          0 2025-10-07 16:12  
/data/TFWCoutput/_SUCCESS  
-rw-r--r--    1 administrator supergroup        32 2025-10-07 16:12  
/data/TFWCoutput/part-r-00000
```

- administrator@administrator:~\$ hadoop fs -ls
/data/TFWCoutput/part-r-00000

```
-rw-r--r-- 1 administrator supergroup          32 2025-10-07 16:12
/data/TFWCoutput/part-r-00000
```

- administrator@administrator:~\$ hadoop fs -cat /data/TFWCoutput/part-r-00000

```
Hadoop2
fast 1
is 2
powerful 1
```

```
administrator@administrator:~$ hadoop fs -ls /data/
Found 2 items
drwxr-xr-x - administrator supergroup          0 2025-10-07 16:08 /data/TFWC
drwxr-xr-x - administrator supergroup          0 2025-10-07 16:12 /data/TFWCoutput
administrator@administrator:~$ hadoop fs -ls /data/TFWCoutput
Found 2 items
-rw-r--r-- 1 administrator supergroup          0 2025-10-07 16:12 /data/TFWCoutput/_SUCCESS
-rw-r--r-- 1 administrator supergroup          32 2025-10-07 16:12 /data/TFWCoutput/part-r-00000
administrator@administrator:~$ hadoop fs -ls /data/TFWCoutput/part-r-00000
-rw-r--r-- 1 administrator supergroup          32 2025-10-07 16:12 /data/TFWCoutput/part-r-00000
administrator@administrator:~$ hadoop fs -cat /data/TFWCoutput/part-r-00000
Hadoop 2
fast 1
is 2
powerful 1
administrator@administrator:~$ █
```